

Lecture 3: Regression and Smoothing

Introduction to Time Series, Fall 2023

Ryan Tibshirani

Related reading: Chapters 2 of Shumway and Stoffer (SS); Chapter 7 of Hyndman and Athanasopoulos (HA).

1 Simple regression

1.1 Population version

- We'll start off by learning the very basics of linear regression, assuming you have not seen it before. A lot of what we'll learn here is not necessarily specific to the time series setting, though of course (especially as the lecture goes on) we'll emphasize the time series angle as appropriate
- A *simple linear regression* model for a response variable y and predictor (or covariate, or feature) variable x is one in which we seek *coefficients* (or parameters) β_0 and β_1 , such that, informally,

$$y \approx \beta_0 + \beta_1 x$$

To be clear, here x, y are all real-valued (rather than multivariate) random variables

- If we had access to the full distributions of x, y , which is what we call the “population version” of regression, then we could ask: what is the best choice of parameters β_0, β_1 with respect to expected squared error?
- Mathematically, we are looking to solve

$$\min_{\beta_0, \beta_1} \mathbb{E}[(y - \beta_0 - \beta_1 x)^2] \tag{1}$$

or in other words, we are asking for the “line of best fit” at the population level. You'll often also hear this referred to as the “least squares” problem

- We can find the answer by differentiating the loss $Q = \mathbb{E}[(y - \beta_0 - \beta_1 x)^2]$ in (1) with respect to each parameter and setting it equal to zero. Differentiating inside the expectation gives:

$$\begin{aligned} \frac{\partial Q}{\partial \beta_0} &= \mathbb{E}[2(\beta_0 + \beta_1 x - y)] = 0 \\ \frac{\partial Q}{\partial \beta_1} &= \mathbb{E}[2x(\beta_0 + \beta_1 x - y)] = 0 \end{aligned}$$

- As you'll show on the homework, solving this pair of equations gives the *population regression coefficients*:

$$\beta_1^* = \frac{\text{Cov}(x, y)}{\text{Var}(x)}, \quad \beta_0^* = \mathbb{E}(y) - \beta_1^* \mathbb{E}(x) \tag{2}$$

- Recalling that $\text{Cor}(x, y) = \text{Cov}(x, y) / \sqrt{\text{Var}(x) \text{Var}(y)}$, we may rewrite the slope as

$$\beta_1^* = \text{Cor}(x, y) \sqrt{\frac{\text{Var}(y)}{\text{Var}(x)}},$$

which shows that it treats x, y *asymmetrically*. This is important to remember. In general, when y is the response and x is the predictor, we speak this relationship as the “regression of y on x ”

1.2 Sample version

- For the “sample version” of linear regression, we seek β_0, β_1 such that for given samples x_i, y_i (covariate and response pairs), $i = 1, \dots, n$,

$$y_i \approx \beta_0 + \beta_1 x_i, \quad i = 1, \dots, n$$

but without access to the full distributions of each x_i and y_i , just given these samples

- We can imagine two ways to proceed:
 1. Start from the population-level formula (2), and use plug-in estimates for the covariance and variance
 2. Start from the population-least least squares problem (1), write down a sample version, then solve it
- Somewhat remarkably, these two strategies end up at the same answer (which need not be the case)
- For strategy 1, we use the sample covariance and sample variance,

$$\widehat{\text{Cov}}(x, y) = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}), \quad \widehat{\text{Var}}(x) = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

where $\bar{x} = \frac{1}{n} \sum_{t=1}^n x_t$ and $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$ are the sample means, and plug these into (2) to get:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}, \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} \quad (3)$$

which we call the *sample regression coefficients*

- For strategy 2, we write down the sample least squares problem

$$\min_{\beta_0, \beta_1} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 \quad (4)$$

- Similar to before, denote the loss in (1) by $Q = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$ and differentiate with respect to each parameter and set it equal to zero:

$$\begin{aligned} \frac{\partial Q}{\partial \beta_0} &= \sum_{i=1}^n 2(\beta_0 + \beta_1 x_i - y_i) = 0 \\ \frac{\partial Q}{\partial \beta_1} &= \sum_{i=1}^n 2x_i(\beta_0 + \beta_1 x_i - y_i) = 0 \end{aligned}$$

- You’ll show on the homework that solving this pair of equations leads you right back to (3)
- The `lm()` function in R performs linear regression. The notation you use is `lm(y ~ x)`, where `y ~ x` is called a “formula”. This can be read as an instruction: “regress y on x ”
- Figure 1 gives an example where we regress chicken prices—our response, y , on time—our predictor x . Just to give you a clear sense, the data are

$$\begin{aligned} y_1 &= 65.58, y_2 = 66.48, y_3 = 65.70, \dots \\ x_1 &= 2001.583, x_2 = 2001.667, x_3 = 2001.750, \dots \end{aligned}$$

where we interpret each value of x as a given year plus a fraction, representing the month of the year

- After running `lm()`, the resulting object is a (special) list, with a lot of useful components. Calling `coef()` on the object gives the regression coefficients



Figure 1: *Linear regression of chicken prices and on time (from SS).*

- Figure 2 gives another example, of a different flavor. Now the response y is itself one time series: cardiovascular mortality in Los Angeles over a certain time period, and the covariate x is itself another time series: particulate levels in Los Angeles over the same time period. The top panel in Figure 2 plots them individually as time series, and the bottom panel plots them together, as a scatter plot, together with the fitted line from linear regression. We can imagine, in a future period (beyond the end date of these time series), using the estimated regression coefficients to predict mortality from particulate levels

1.3 Prediction: ex-ante and ex-post

- We can use the estimated coefficients $\hat{\beta}_0, \hat{\beta}_1$ in (3) from linear regression estimates to make *predictions* about the response given a new predictor value x_{new} . This prediction is

$$\hat{y}_{\text{new}} = \hat{\beta}_0 + \hat{\beta}_1 x_{\text{new}},$$

where the “hat” notation on the left-hand side emphasizes that it is an not observed, but an estimated (predicted) value of the response

- In time series context, we often will use the term *forecasting* synonymously with *prediction*. In the time series, there is an interesting distinction between two types of forecasts, based on *whether or not the predictor value x_{new} needed to make forecasts is available in advance*. Specifically:
 - An *ex-ante forecast* is a “true” forecast, using only information that is available at the time the forecast was issued. So the predictor values need to either be available, or themselves be forecasted. For instance, we can make ex-ante forecasts in the chicken regression example
 - An *ex-post forecast* is made using later information on the predictors. So we wait until x_{new} is observed, then issue our forecast. For instance, we can make ex-post forecasts in the mortality



Figure 2: *Linear regression of cardiovascular mortality on particulate levels in Los Angeles (from SS).*

regression example (but cannot easily make ex-ante forecasts unless we somehow could forecast particular levels into the future)

- One way around the circumvent the potential difficulty of making ex-ante forecasts is to use *lagged predictors*; we'll discuss this and other forecasting issues in more detail later

1.4 A note on assumptions and philosophy

- What have we assumed above? *Nothing*. That is, *we do not need to assume that the true relationship between y and x is linear in order to perform linear regression as in (3)*
- Thus, in general, there need not be any “true regression coefficients” that we’re actually tracking ... but, we can always think of the sample estimates $\hat{\beta}_0, \hat{\beta}_1$ in (3) as tracking the population quantities β_0^*, β_1^* in (2). The latter are basically also always well-defined, regardless of linearity. Recall, they are the viewed as the best linear approximation at the population level
- So, to be clear, we can always fit sample coefficients $\hat{\beta}_0, \hat{\beta}_1$ and use them to make predictions (forecasting, in time series). Sometimes we call this our “working model”: to use a linear working model is a modeling decision, not an assumption
- If this predicts well (has good accuracy), then our working model was a good decision, and depending on our use case, we may not even care about whether the true model is linear (or related assumptions in classical linear modeling)
- Meanwhile, for use cases would that require *inference*, we require lots of assumptions. More on this, later

2 Multiple regression

2.1 Population version

- What about the case where we have more than out covariate? This is called *multiple linear regression*. Now let $x = (x_1, \dots, x_p) \in \mathbb{R}^p$ be a random vector of covariates, with each entry x_j being an individual covariate of interest
- We seek β_0 , which is an intercept term as before, and also a whole coefficient vector $\beta = (x_1, \dots, x_p) \in \mathbb{R}^p$, such that

$$y \approx \beta_0 + \underbrace{\beta_1 x_1 + \dots + \beta_p x_p}_{\beta^\top x}$$

- Our convention throughout this class will be to *treat all vectors as column vectors*. Thus a^\top , which is the transpose of a vector a , is a row vector, and for vectors $a, b \in \mathbb{R}^d$, we can use $a^\top b = a_1 b_1 + \dots + a_d b_d$ to denote their inner product.
- (Note that of course $a^\top b = b^\top a$, so it doesn't matter whether we write $\beta^\top x$ or $x^\top \beta$ in our model)
- As in (1), we define the population-level regression coefficients by minimizing expected squared error,

$$\min_{\beta_0, \beta} \mathbb{E}[(y - \beta_0 - x^\top \beta)^2] \quad (5)$$

- The solution, which you can think of as generalizing (2), is

$$\beta^* = \text{Cov}(x)^{-1} \text{Cov}(x, y), \quad \beta_0^* = \mathbb{E}(y) - \mathbb{E}(x)^\top \beta^* \quad (6)$$

- Let's check that the dimensions make sense: here $\text{Cov}(x) \in \mathbb{R}^{p \times p}$, a $p \times p$ matrix of real values; the element in its i^{th} row and j^{th} column is

$$[\text{Cov}(x)]_{ij} = \text{Cov}(x_i, x_j)$$

Also $\text{Cov}(x, y) \in \mathbb{R}^p$, a p -dimensional vector, with i^{th} entry

$$[\text{Cov}(x, y)]_i = \text{Cov}(x_i, y)$$

So $\text{Cov}(x)^{-1} \text{Cov}(x, y) \in \mathbb{R}^p$, itself a p -dimensional vector, which is what we need for β^* in (6). Similarly, you can check that the dimensions make sense for β_0^*

- To derive (6) as the minimizer in (5), we can again differentiate with respect to each β_j and set the result to zero, but the calculation is a little more difficult (maybe it'll be a bonus on the homework)

2.2 Sample version

- The sample version of multiple linear regression falls out of the population version entirely analogously, as it did in the simple linear regression case. We seek $\beta_0 \in \mathbb{R}$ and $\beta_1 \in \mathbb{R}^p$ such that for given samples x_i, y_i (covariate and response pairs), $i = 1, \dots, n$,

$$y_i \approx \beta_0 + \underbrace{\beta_1 x_{i1} + \dots + \beta_p x_{ip}}_{x_i^\top \beta}, \quad i = 1, \dots, n$$

- We can do this either by plug-in estimates from (6), or by writing down a sample version of the least squares problem (5) and solving it, and again they will both lead to the same answer
- Let's pursue the latter. It will be convenient to adopt the following convention, which alleviates us from keeping track of an explicit intercept term β_0 , without any loss of generality. We simply write

$$y_i \approx x_i^\top \beta, \quad i = 1, \dots, n$$

without intercept. Then all results that we will derive can be translated to the model with intercept via the following trick: we redefine each vector x_i so that it has a 1 prepended to it:

$$x = (1, x_1, \dots, x_p)$$

Then we read off results with this new transformation: post-transformation, the first entry of β serves as the intercept, and the rest serve as the coefficients multiplying each x_j

- (There is another way to get rid of the intercept as well, which we'll learn when we connect multiple to simple linear regression a bit later, which may be more intuitive to some of you)
- The sample least squares problem for multiple regression is now as follows:

$$\min_{\beta \in \mathbb{R}^p} \sum_{i=1}^n (y_i - x_i^\top \beta)^2 \tag{7}$$

- The solution, obtained again by taking derivatives and setting equal to zero, is

$$\hat{\beta} = \left(\sum_{i=1}^n x_i x_i^\top \right)^{-1} \sum_{i=1}^n x_i y_i \tag{8}$$

- You can check that the dimensions all make sense here (that the right-hand side in (8) produces a p -dimensional vector)

2.3 Matrix notation

- It is more convenient, once you've sufficiently familiarized yourself with matrix notation, to recast regression in terms of matrices and vectors. Let $y = (y_1, \dots, y_n) \in \mathbb{R}^n$ be the vector of our response values, and $X \in \mathbb{R}^{n \times p}$ the matrix of our predictor vectors, whose i^{th} row is x_i^\top

- In other words,

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}, \quad X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

- Recall, for a coefficient vector $\beta \in \mathbb{R}^p$, the matrix-vector product $X\beta \in \mathbb{R}^n$, which to emphasize is an n -dimensional vector, is

$$X\beta = \begin{bmatrix} x_1^\top \beta \\ x_2^\top \beta \\ \vdots \\ x_n^\top \beta \end{bmatrix}$$

This means that we can write our sample working model compactly as

$$y \approx X\beta$$

- Recall, the Euclidean norm $\|\cdot\|$ of a vector $a \in \mathbb{R}^d$ is defined as $\|a\|^2 = \sum_{i=1}^d a_i^2$. This means we can write our sample least squares problem (7) compactly as

$$\min_{\beta \in \mathbb{R}^p} \|y - X\beta\|^2 \quad (9)$$

- Finally, the sample least squares estimates (8) can be written as

$$\hat{\beta} = (X^\top X)^{-1} X^\top y \quad (10)$$

This form (10) is by far the more commonly-used (and easily-remembered) form for the least squares coefficient estimates, compared to (8). You'll prove the equivalence between the two forms (8), (10) on the homework

- Important technical note: above in (10) (also in (8)), we have implicitly assumed that the features—the columns of X —are *linearly independent*. This *can only happen if* $p \leq n$, i.e., if we have no more features than samples. Otherwise, $X^\top X$ will not have an inverse, strictly speaking. This is not the end of the world and there are many interesting things to talk about (along the lines of regularization) when $p > n$, but it's beyond our scope right now

2.4 Multiple vs simple: a connection

- There is quite an interesting connection between multiple regression of y on x and simple regression of y on each x_j , the latter often referred to as a *marginal linear regression*

****Warning** **Warning** **Warning****

There will be a notational clash between x_i as used above and x_j as will be used here. Previously, we used $x_i \in \mathbb{R}^p$ to refer to vector containing all feature values for the i^{th} sample. Here, we are going to use $x_j \in \mathbb{R}^n$ to refer to the vector containing the j^{th} feature values measured over all samples.

For example, suppose we have two features: apples and bananas, and we measure the quantity of each across $n = 100$ households. Then the previous subsection used $x_i \in \mathbb{R}^2$ as the number of apples and bananas in the i^{th} household. But here we'll use $x_j \in \mathbb{R}^{100}$ as the number of apples (if $j = 1$) in the 100 households, or bananas (if $j = 2$) in the households. Get it?

Since the beginning of ~~time~~ regression, scholars have run up against this problem and have racked their brains for notational solutions. But there is no real good notational solution to this. (Yes there are lots of options but all of them have some ugliness to them.)

We'll just make to use $x_i \in \mathbb{R}^p$ to always refer to all features for the i^{th} sample, and $x_j \in \mathbb{R}^n$ to always refer to the j^{th} feature for all samples (and try to never mix indices). In other words, the i^{th} row and j^{th} column of the feature matrix X , respectively. So you just have to remember that i indexes samples (rows), and j indexes features (columns).

- Now that we've gotten that important notational piece out of the way, we will define more quantities in order to describe the connection between multiple and marginal regression
- Fix any j (arbitrary). Given a single feature $x_j = (x_{1j}, \dots, x_{nj}) \in \mathbb{R}^n$, and response vector $y = (y_1, \dots, y_n) \in \mathbb{R}^n$, the marginal (or simple) regression of y on x_j , *without intercept*, is a very similar formula to what you saw in (3):

$$\tilde{\beta}_j = \frac{\sum_{i=1}^n x_{ij} y_i}{\sum_{i=1}^n x_{ij}^2}$$

(Note: the effect of the intercept is only that it centers the values of x_{ij} around their sample average, which you'll revisit on the homework)

- We can rewrite this succinctly in inner product notation as:

$$\tilde{\beta}_j = \frac{x_j^\top y}{x_j^\top x_j} \quad (11)$$

- Meanwhile, let's consider the j^{th} estimated regression coefficient $\hat{\beta}_j$ from the multiple regression of y on X , as in (10). Is this the same? That is, is the j^{th} component of (10) the same as (11)?
- The answer is generally no! Putting in other features into the linear regression will generally affect the estimated coefficient for x_j , i.e., will alter its “predictive influence” on y
- However, there is a precise connection between multiple regression and marginal regression. Let's write $X_{-j} \in \mathbb{R}^{n \times p}$ for the feature matrix but after dropping $x_j \in \mathbb{R}^n$, the j^{th} column. Suppose we:

- Regress y on X_{-j} (a regression of y on all but the j^{th} feature), yielding an estimated coefficient vector $\hat{\alpha} \in \mathbb{R}^{p-1}$, and residual

$$\hat{y}^{-j} = y - X_{-j} \hat{\alpha} \quad (12)$$

- Regress x_j on X_{-j} (a regression of x_j on all of the other features), yielding an estimated coefficient vector $\hat{\theta} \in \mathbb{R}^{p-1}$, and residual

$$\hat{x}_j^{-j} = x_j - X_{-j} \hat{\theta} \quad (13)$$

- In these residuals (12), (13), we have “regressed out the influence” of all other features on each of y and x_j . Then, after removing such influence, if we perform a marginal regression of \hat{y}^{-j} on \hat{x}_j^{-j} , we get precisely the j^{th} multiple regression coefficient:

$$\hat{\beta}_j = \frac{(\hat{x}_j^{-j})^\top \hat{y}^{-j}}{(\hat{x}_j^{-j})^\top \hat{x}_j^{-j}} \quad (14)$$

- In other words, (14) connects multiple regression to marginal regression: it shows that the j^{th} coefficient in a multiple regression (10) is equivalent to a marginal regression of y on x_j , but only after we have accounted for the effects of all the other predictors, by “regressing them out”
- The best way to understand the relationship between (14) and (10) is geometrically (which is also a great way to view linear regression in general) but that perspective requires a bit more advanced linear algebra, which we won't cover. For now, you can just think of the following: the bigger the *correlations* between x_j and columns of X_{-j} , the bigger the effect will be in (13), where we regress out X_{-j} from x_j , and this will make the multiple (14) and marginal (11) coefficients quite different from each other

- On the other hand, when x_j is uncorrelated with each column of X_{-j} , then the residual \hat{x}_j^{-j} in (13) is no different from x_j , and in fact one can show that the multiple regression coefficient in (14) and the marginal one in (11) are exactly the same
- Figure 3 revisits the cardiovascular mortality example from earlier. Now we perform the regression of cardiovascular mortality on two features: particulate levels and temperature. (In R we can regress y on two features $x1$ and $x2$ by running `lm(y ~ x1 + x2)`.) For each feature, the coefficients from multiple regression aren't too different from the marginal regression coefficients (as is seen in the bottom panel by the slopes of the lines), though the intercept changes noticeably. The relatively small change in the coefficients on particulate levels and temperature is due to the fact that these two are not very correlated (as is seen by looking at their time series, which are a bit “out of phase” with each other)

2.5 Interlude: best linear unbiased estimator

- Briefly, we discuss an important optimality property here of least squares estimates, before moving on to classical inferential results next. We are going to need to introduce an assumption for this part: we assume that the response vector $y \in \mathbb{R}^n$ is related to the feature matrix $X \in \mathbb{R}^{n \times p}$ by

$$y = X\beta + \epsilon, \quad \text{where } \epsilon \sim \text{WN}(0, \sigma^2 I) \quad (15)$$

for some unknown coefficient vector $\beta \in \mathbb{R}^p$, and a white noise vector $\epsilon \in \mathbb{R}^n$

- In (15), $\epsilon \sim \text{WN}(0, \sigma^2 I)$ is our notation for a white noise vector: the first argument specifies that the mean is zero, $\mathbb{E}(\epsilon) = 0$, and the second argument specifies that the covariance satisfies $\text{Cov}(\epsilon) = \sigma^2 I$, where I is the $n \times n$ identity matrix. That is, the components satisfy $\text{Var}(\epsilon_i) = \sigma^2$ and $\text{Cov}(\epsilon_i, \epsilon_j) = 0$ whenever $i \neq j$
- Importantly, in (15), the feature matrix X is assumed to be *fixed* (not random). Therefore, we can also write (15) even more compactly as

$$y \sim \text{WN}(X\beta, \sigma^2 I)$$

This is a fairly strong assumption. We'll go into more why in the next section (when we go even further and assume normality of the error distribution), but for now we'll just emphasize that we are assuming that the mean of the response is truly linear in the features, and that the errors are homoskedastic—they have equal variance regardless of the feature values

- Ok! So under the model (15), what can we say about the least squares estimates in (10)? First, note that these are *unbiased* for the true coefficients β :

$$\begin{aligned} \mathbb{E}(\hat{\beta}) &= \mathbb{E}[(X^\top X)^{-1} X^\top y] \\ &= (X^\top X)^{-1} X^\top \mathbb{E}(y) \\ &= (X^\top X)^{-1} X^\top X\beta \\ &= \beta \end{aligned}$$

- This implies it is unbiased for any contrast of β . This is the term we sometimes give to an estimand of the form $a^\top \beta$, for an arbitrary vector $a \in \mathbb{R}^d$. Observe,

$$\begin{aligned} \mathbb{E}(a^\top \hat{\beta}) &= a^\top \mathbb{E}(\hat{\beta}) [(X^\top X)^{-1} X^\top y] \\ &= (X^\top X)^{-1} X^\top \mathbb{E}(y) \\ &= (X^\top X)^{-1} X^\top X\beta \\ &= \beta \end{aligned}$$

- One common way to measure the quality of an estimator is its *mean squared error* (MSE), which for the least squares estimator $a^\top \hat{\beta}$ of the contrast $a^\top \beta$, is

$$\text{MSE}(a^\top \hat{\beta}) = \mathbb{E}[(a^\top \hat{\beta} - a^\top \beta)]^2$$

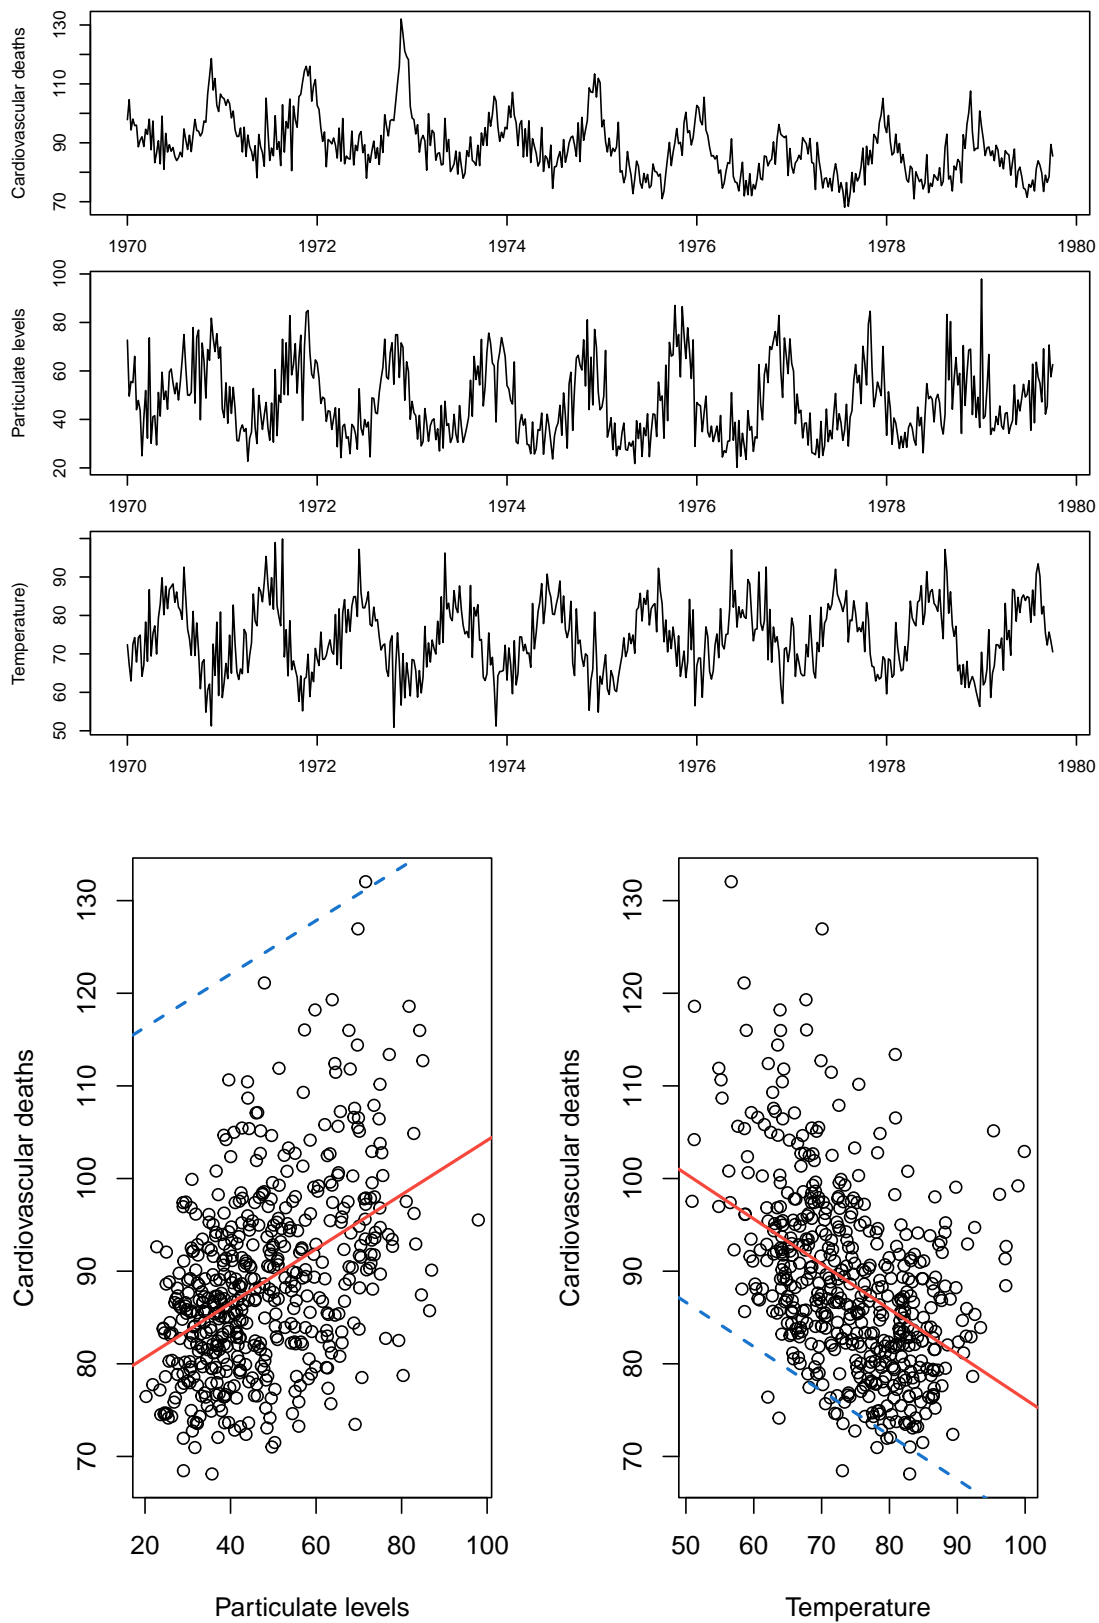


Figure 3: *Linear regression of cardiovascular mortality on particulate levels and temperature (this is multiple regression, with two features) in Los Angeles (from SS). The solid red lines denote the estimates from marginal regression; the dashed blue from multiple regression.*

- To be clear, the expectation here is being taken with respect to data from the model (15)
- With respect to MSE, how good is the least squares estimator? It is, in a certain precise sense, the *best*. It is both unbiased for $a^\top \beta$, as proved above, and also a *linear* estimator as a function of y : we can write

$$a^\top \hat{\beta} = \underbrace{(a^\top X^\top X)^{-1} X^\top}_{b^\top} y$$

That is, $\hat{\beta} = b^\top y$, where $b = X(X^\top X)^{-1}a$

- Note: linearity of the estimator has nothing to do with linearity of the true regression model! The former (linearity of the estimator) is a statement about linearity in y , the response; the latter (linearity of the true regression model) is a statement about linearity in X , the features. There is a notational collision, but linearity is really referring to different things here
- Now for the main event: the *Gauss-Markov theorem* tells us that the least squares estimator is the best linear unbiased estimator (BLUE) of $a^\top \beta$. In other words, for any other linear estimator $c^\top y$, such that $\mathbb{E}(c^\top y) = a^\top \beta$ (unbiasedness), we have

$$\text{MSE}(a^\top \hat{\beta}) \leq \text{MSE}(c^\top y)$$

- A proof of this fact follows from the geometric perspective on least squares, which we won't cover, but you can ask about it in office hours if you are curious
- (An interesting side note! Econometricians have been recently arguing about whether or not we can drop the “L” from BLUE: is least squares the BUE? That is, is least squares the best unbiased estimator, period—best among all unbiased estimators, not just linear ones? The answer is ... yes, in a sense, but it depends on how you set up the problem, and in certain problem settings the only unbiased estimators are linear in y anyway.¹)

3 Classical inference

3.1 Here comes the assumptions

4 Diagnostics

5 Forecasting

6 Smoothing

¹See Hansen (2022), “A modern Gauss-Markov theorem”, and then Pötscher and Preinerstorfer (2022), “A modern Gauss-Markov theorem? Really?”, and Portnoy (2022), “Linearity of unbiased linear model estimators” (2022). A nice, gentle overview is given by Allison (2022), “Is OLS BLUE or BUE?”, <https://statisticalhorizons.com/is-ols-blue-or-bue/>. A masterful, but much more mathematical overview is given in Lei and Wooldrige (2022), who also weave in important historical results that have been overlooked, and prove new ones as well.